


| | |
|-----------------------|----------------|
| Name | Shubham Golwal |
| UID no. | 2020300015 |
| Experiment No. | 01 |

| | |
|----------------------------|---|
| AIM: | Concurrency control |
| Program 1 | |
| Problem Statement : | Perform concurrency control in PostgreSQL Database with college database. |
| Theory : | <p>What are transactions? A transaction is a very small unit of a program and it may contain several lowlevel tasks. A transaction in a database system must maintain Atomicity, Consistency, Isolation, and Durability – commonly known as ACID properties – in order to ensure accuracy, completeness, and data integrity.</p> <p>What is concurrency control? Concurrency control concept comes under the Transaction in database management system (DBMS). It is a procedure in DBMS which helps us for the management of two simultaneous processes to execute without conflicts between each other, these conflicts occur in multi user systems. Concurrency can simply be said to be executing multiple transactions at a time. It is required to increase time efficiency. If many transactions try to access the same data, then inconsistency arises. Concurrency control required to maintain consistency data. For example, if we take ATM machines and do not use concurrency, multiple persons cannot draw money at a time in different places. This is where we need concurrency.</p> <p>The advantages of concurrency control are as follows –</p> <ul style="list-style-type: none"> • Waiting time will be decreased. • Response time will decrease. • Resource utilization will increase. • System performance & Efficiency is increased. |



All the descendents release the held locks and retained locks

**credits to gfg*

Output

Drop:

The screenshot shows the pgAdmin 4 interface. The SQL editor contains the query: `DROP TABLE IF EXISTS college;`. The Messages tab at the bottom shows the output: `DROP TABLE` and `Query returned successfully in 52 msec.`

Create:

The screenshot shows the pgAdmin 4 interface. The left sidebar displays the database structure: Servers > PostgreSQL > Database > college. The main SQL editor window is titled 'Properties ADBMS Lab/postgres@PostgreSQL 15*'. It contains the following SQL code:

```
1  -- DROP TABLE IF EXISTS Laptops;
2
3
4  CREATE TABLE college (
5      id INT GENERATED BY DEFAULT AS IDENTITY,
6      name VARCHAR(100) NOT NULL,
7      age DEC(15,2) NOT NULL,
8      subject VARCHAR(100) NOT NULL,
9      score DEC(15,2) NOT NULL,
10     PRIMARY KEY(id)
11 );
12
13
```

Below the SQL editor, the 'Messages' tab is active, showing the message: 'CREATE TABLE' and 'Query returned successfully in 132 msec.'

Insert:

The screenshot shows the pgAdmin 4 interface. The left sidebar displays the database structure: Servers > PostgreSQL > Database > college. The main SQL editor window is titled 'Properties ADBMS Lab/postgres@PostgreSQL 15*'. It contains the following SQL code:

```
7  -- age DEC(15,2) NOT NULL,
8  -- subject VARCHAR(100) NOT NULL,
9  -- score DEC(15,2) NOT NULL,
10 -- PRIMARY KEY(id)
11 -- );
12
13
14
15 INSERT INTO college(name, age, subject, score)
16 VALUES('Shubham', 20 , 'ADBMS', 85);
17
18
19
```

Below the SQL editor, the 'Messages' tab is active, showing the message: 'INSERT 0 1' and 'Query returned successfully in 76 msec.'

Query Query History

```
31
32
33 SELECT
34     id,
35     name,
36     age,
37     subject,
```

Data Output Messages Explain × Notifications




| | id [PK] integer | name character varying (100) | age numeric (15,2) | subject character varying (100) | score numeric (15,2) |
|---|--------------------|---------------------------------|-----------------------|------------------------------------|-------------------------|
| 1 | 1 | Shubham | 20.00 | ADBMS | 85.00 |

Initial insert into database

Begin Transaction:

Help

SQL Properties ADBMS Lab/postgres@PostgreSQL 15*



Query Query History

```
17
18
19
20 BEGIN TRANSACTION;
21
22 INSERT INTO college(name, age, subject, score)
23 VALUES('Varun', 19 , 'DBMS', 75);
24
25 INSERT INTO college(name, age, subject, score)
26 VALUES('Noman', 21 , 'DSA', 95);
27 INSERT INTO college(name, age, subject, score)
28 VALUES('Suyog', 20 , 'SE', 65);
29 INSERT INTO college(name, age, subject, score)
30 VALUES('Pratik', 21 , 'TOC', 86);
31
32
33 SELECT
34     id,
```

Data Output Messages Explain × Notifications

INSERT 0 1

Query returned successfully in 85 msec.

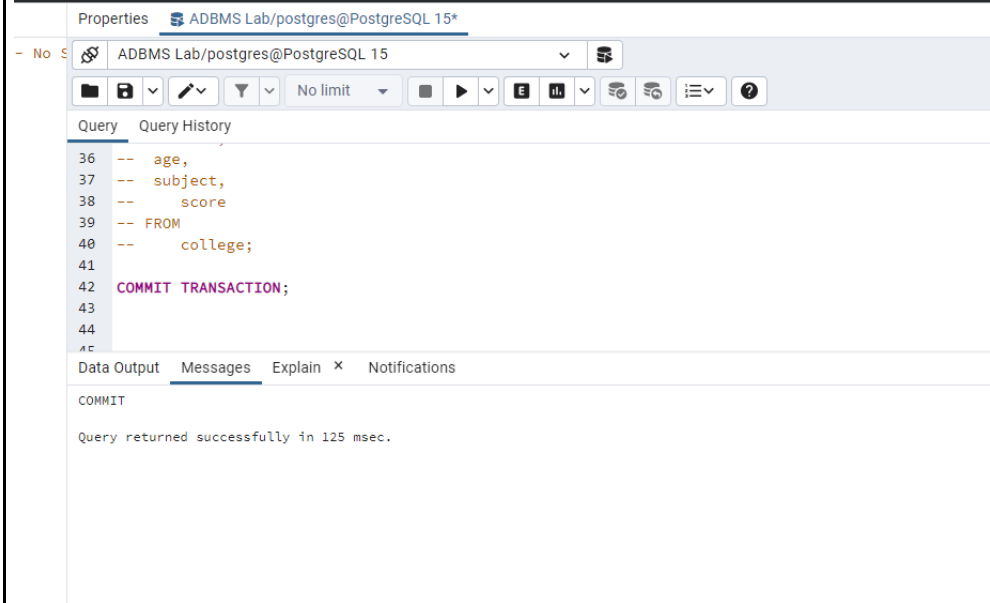
Data Output Messages Explain × Notifications



| | id [PK] integer | name character varying (100) | age numeric (15,2) | subject character varying (100) | score numeric (15,2) |
|---|--------------------|---------------------------------|-----------------------|------------------------------------|-------------------------|
| 1 | 1 | Shubham | 20.00 | ADBMS | 85.00 |
| 2 | 2 | Varun | 19.00 | DBMS | 75.00 |
| 3 | 3 | Noman | 21.00 | DSA | 95.00 |
| 4 | 4 | Suyog | 20.00 | SE | 65.00 |
| 5 | 5 | Pratik | 21.00 | TOC | 86.00 |

Started a new transaction and insert a new student into the college table:

Commit:



The screenshot shows a PostgreSQL query editor interface. The top bar indicates the connection is to 'ADBMS Lab/postgres@PostgreSQL 15*'. Below the toolbar, the 'Query' tab is active, displaying a SQL script with line numbers 36 to 44. The script contains a multi-line INSERT statement followed by a COMMIT TRANSACTION statement on line 42. The 'Messages' tab is also visible, showing a 'COMMIT' message and a status report: 'Query returned successfully in 125 msec.'

```
36 -- age,  
37 -- subject,  
38 -- score  
39 -- FROM  
40 -- college;  
41  
42 COMMIT TRANSACTION;  
43  
44  
45
```

COMMIT
Query returned successfully in 125 msec.

Made the changes become visible to other sessions (or users), committed the transaction by using the COMMIT TRANSACTION statement

Alternative: COMMIT WORK, COMMIT;

After executing the COMMIT statement, PostgreSQL also guarantees that the change will be durable if a crash happens.

-----X-----

Update:

Currently score at id = 2 is 75.

Properties ADBMS Lab/postgres@PostgreSQL 15*

No S ADBMS Lab/postgres@PostgreSQL 15

Query Query History

```

56 -- UPDATE
57 BEGIN;
58
59 UPDATE college
60 SET score = (score/10) + 40
61 WHERE id = 2;
62
63 SELECT
64     id,
65     name,
66     age,
67     subject,
68     score
69 FROM
70     college;
71
72 COMMIT;

```

63 SELECT
64 id,

Data Output Messages Explain × Notifications

| | id [PK] integer | name character varying (100) | age numeric (15,2) | subject character varying (100) | score numeric (15,2) |
|---|--------------------|---------------------------------|-----------------------|------------------------------------|-------------------------|
| 1 | 1 | Shubham | 20.00 | ADBMS | 85.00 |
| 2 | 3 | Noman | 21.00 | DSA | 95.00 |
| 3 | 4 | Suyog | 20.00 | SE | 65.00 |
| 4 | 5 | Pratik | 21.00 | TOC | 86.00 |
| 5 | 2 | Varun | 19.00 | DBMS | 47.50 |

Changes is made to id =2 [Varun] and score is updated to 47.5
After committing changes are made visible.

Rollback:

Updated but not committed:

```

56 -- UPDATE
57 BEGIN;
58
59 UPDATE college
60 SET score = (score/10) + 40
61 WHERE id = 2;
62
63 -- SELECT
64 --     id,

```

Score at id = 2 is 47.5.

```

83
84 ROLLBACK TRANSACTION;

```

| | id [PK] integer | name character varying (100) | score numeric (15,2) | age numeric (15,2) |
|---|--------------------|---------------------------------|-------------------------|-----------------------|
| 1 | 1 | Shubham | 85.00 | 20.00 |
| 2 | 2 | Varun | 75.00 | 19.00 |
| 3 | 3 | Noman | 95.00 | 21.00 |
| 4 | 4 | Suyog | 65.00 | 20.00 |
| 5 | 5 | Pratik | 86.00 | 21.00 |

After executing rollback score at id = 2 is reverted to 75

But if transaction is committed then we can't revert back.

```

82 -- ROLLBACK
83
84 ROLLBACK TRANSACTION;

```

| Data Output | Messages | Explain | × | Notifications |
|--|----------|---------|---|---------------|
| <p>WARNING: there is no transaction in progress</p> <p>ROLLBACK</p> <p>Query returned successfully in 57 msec.</p> | | | | |

Serializability:

I'm gonna start 2 new transactions, then set their isolation level to serializable.

Properties ADBMS Lab/po... ADBMS Lab/postgres@PostgreSQL 15*

ADBMS Lab/postgres@PostgreSQL 15

Query Query History

```
1 -- SELECT * from college;  
2 set transaction isolation level serializable;  
3 show transaction isolation level;  
4  
5
```

Data Output Messages Notifications

| | transaction_isolation text |
|---|-------------------------------|
| 1 | serializable |

Get sum of all the scores of students

Properties ADBMS Lab/po... ADBMS Lab/postgres@PostgreSQL 15*

ADBMS Lab/postgres@PostgreSQL 15

Query Query History

```
1 -- SELECT * from college;  
2 -- set transaction isolation level serializable;  
3 -- show transaction isolation level;  
4  
5 select sum(score) from college;
```

Data Output Messages Notifications

| | sum numeric |
|---|----------------|
| 1 | 378.50 |

Properties ADBMS Lab/po... ADBMS Lab/postgres@PostgreSQL 15*

No S ADBMS Lab/postgres@PostgreSQL 15

Query Query History

```

1 -- SELECT * from college;
2 -- set transaction isolation level serializable;
3 -- show transaction isolation level;
4
5 -- select sum(score) from college;
6
7 insert into college(name, age, subject, score) values ('Mikey', 19, 'BEE', 46) re

```

Data Output Messages Notifications

| | id [PK] integer | name character varying (100) | age numeric (15,2) | subject character varying (100) | score numeric (15,2) |
|---|--------------------|---------------------------------|-----------------------|------------------------------------|-------------------------|
| 1 | 6 | Mikey | 19.00 | BEE | 46.00 |

Query Query History

```

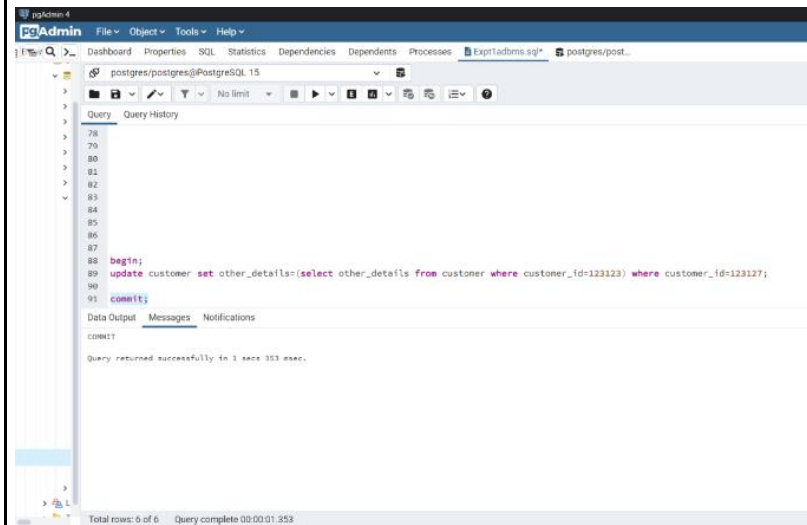
1 SELECT * from college;
2 -- set transaction isolation level serializable;
3 -- show transaction isolation level;
4
5 -- select sum(score) from college;
6
7 -- insert into college(name, age, subject, score) values ('Mikey', 19, 'BEE', 46) returnf
8
9

```

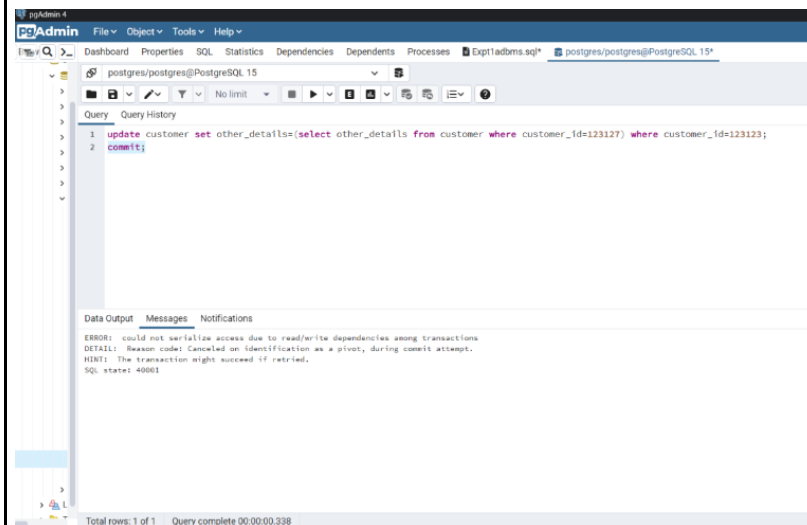
Data Output Messages Notifications

| | id [PK] integer | name character varying (100) | age numeric (15,2) | subject character varying (100) | score numeric (15,2) |
|---|--------------------|---------------------------------|-----------------------|------------------------------------|-------------------------|
| 1 | 1 | Shubham | 20.00 | ADBMS | 85.00 |
| 2 | 3 | Noman | 21.00 | DSA | 95.00 |
| 3 | 4 | Suyog | 20.00 | SE | 65.00 |
| 4 | 5 | Pratik | 21.00 | TOC | 86.00 |
| 5 | 2 | Varun | 19.00 | DBMS | 47.50 |
| 6 | 6 | Mikey | 19.00 | BEE | 46.00 |

Serialized updates:
Window 1:



Window 2:



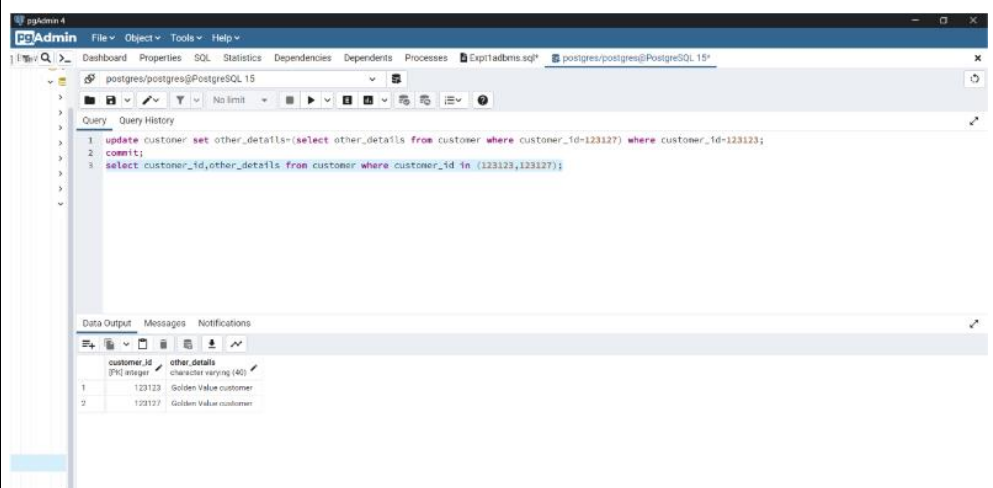
Error: could not serialize access due to read/write dependencies among transactions **DETAIL:**

Reason code: Cancelled on identification as a pivot, during commit attempt.

HINT: The transaction might succeed if retried. SQL state: 40001

Why has it happened?

When we committed first update operation it worked fine. But when we committed second one on same tuple then it throws serializable error.



Checking serializable:

Window 2:

Is this equivalent to any serializable schedule?

Yes, this schedule is conflict serializable. Because we are performing write operation from both transactions at the same time.

Conclusion:

- I learnt meaning of serializability and concurrency in a transaction.
- I also learnt that various Transaction language (TCL) commands are used for managing and controlling the transactions.
- Executed serializability and under the conflict in serializability.